

W, 10/07/19

Fall'19 CSCE 629

Analysis of Algorithms

Fang Song
Texas A&M U


Lecture 15

- Dijkstra's algorithm cont'd
- Interval scheduling

Reflection on Dijkstra: **greedy** stays ahead

- Known region R : in which the shortest distance to s is known
- Growing R : adding v that has the **shortest** distance to s
- How to Identify v ? The one that **minimizes** $d(u) + l(u, v)$ for $u \in R$

Shortest path to some u in known region, followed by a single edge (u, v)



```
Dijk( $G, s$ ) // initialize  $d(s) = 0, d(u) = \infty, R = \emptyset$   
While  $R \neq V$   
    Pick  $v \notin R$  w. smallest  $d(v)$  // by Priority Q  
    Add  $v$  to  $R$   
    For all edges  $(v, w) \in E$   
        If  $d(v) > d(u) + l(u, v)$   
             $d(v) \leftarrow d(u) + l(u, v)$ 
```

Contrast with Bellman-Ford

- Dijkstra (Greedy) $O((m + n) \log n)$

$$d(v) = \min_{u \in R} d(u) + l(u, v)$$

- Positive weight: no need to wait; more edges in a path do not help

- Bellman-Ford (Dynamic programming) $O(mn)$

$$OPT(i, v) = \min \left\{ OPT(i - 1, v), \min_{v \rightarrow w \in E} \{ OPT(i - 1, w) + l_{v \rightarrow w} \} \right\}$$

❖ Global vs. Local

- Dijkstra's requires **global** information: known region & which to add
- Bellman-Ford uses only **local** knowledge of neighbors, suits **distributed** setting

Network routing: distance-vector protocol

■ Communication network

- Nodes: routers
- Edges: direct communication links
- Cost of edge: delay on link.

naturally nonnegative, but
Bellman-Ford used anyway!

■ Distance-vector protocol ["routing by rumor"]

- Each router maintains a vector of shortest-path lengths to every other node (distances) and the first hop on each path (directions).
- Algorithm: each router performs separate computations for each potential destination node.


■ Path-vector protocol: coping with dynamic costs

Correctness of Dijkstra's algorithm

Invariant. For each node $u \in R$, $d(u)$ is the length of a shortest $s - u$ path


Proof. (By induction on size of R)

- Base case: $|R| = 1$ trivial
- Induction hypothesis: true for $|R| = k \geq 1$
 - Let v be the next node added to R and (u, v) be the chosen edge. Call this $s - u - v$ path P .
 - Consider any $s - v$ path Q . [Next show it's no shorter than P]
 - Let (x, y) be the first edge in Q leaving R ; let Q' be the $s - x$ segment
 - $l(Q) \geq l(Q') + l(x, y) \geq d(x) + l(x, y) \geq l(P)$; because Dijkstra's picked v in this iteration (node outside R with shortest distance to s)



Recall: weighted interval scheduling

- **Input.** n jobs; job j starts at s_j , finishes at f_j , weight w_j
- **Output.** Subset of mutually compatible jobs of maximum weight



- DP algorithm $O(n \log n)$

Today

$O(n \log n)$ Greedy algorithm for $w_j = 1$.

Greedy strategies

Recall. DP recurrence.

$\text{OPT}(j)$ = value of optimal solution to jobs $1, 2, \dots, j$

$$\text{OPT}(j) = \begin{cases} 0 & \text{if } j = 0 \\ \max\{\text{OPT}(j-1), w_j + \text{OPT}(\text{pre}(j))\} & \text{otherwise} \end{cases}$$

- **Greedy:** be lazy & pick the next compatible job that “looks nice”
 - **Earliest start time:** ascending order of s_j .
 - **Earliest finish time:** ascending order of f_j .
 - **Shortest interval:** ascending order of $f_j - s_j$.
 - **Fewest conflicts:** the one that conflicts the **least** number of jobs go first.
- **Exercise.** Find counterexamples for each strategy (if possible)

Greedy: counterexamples

☹️ Earliest start time:



☹️ Shortest interval:



☹️ Fewest conflicts:



😊 Earliest finishing time

Greedy Algorithm: earliest finishing time

IntScheduling ($\{s_j, f_j\}$)

1. Sort by finishing time so that $f_1 \leq f_2 \leq \dots \leq f_n \rightarrow O(n \log n)$
2. $A \leftarrow \emptyset$ // set of selected jobs
3. For $j = 1, \dots, n$
 - If j compatible with A
 - $A \leftarrow A \cup \{j\}$


} $n \times O(1)$

- Running time: $O(n \log n)$
- Correctness: proof by contradiction
 - Suppose greedy is not optimal
 - Consider an optimal strategy: one that agrees with Greedy for as many initial jobs as possible
 - Look at the first place that they differ: show a new optimal that agrees with greedy **more**

Greedy Algorithm: correctness

Proof (by contradiction): Suppose greedy is not optimal

- Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy
- Let j_1, j_2, \dots, j_m be set of jobs in the optimal solution OPT where $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the **largest possible** value of r
- Sub i_{r+1} for j_{r+1} in OPT: still feasible and optimal (OPT'); but agrees with Greedy at $r + 1$ positions; contradicts the maximality of r




Interval Partitioning Problem

Scheduling classes

- **Input.** Lectures $\{s_j, f_j\}$
- **Output.** **Minimum number** of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Can you do better? 10 lectures scheduled in **4** classrooms




Interval Partitioning Problem

Scheduling classes

- **Input.** Lectures $\{s_j, f_j\}$
- **Output.** **Minimum number** of classrooms to schedule all lectures so that no two occur at the same time in the same room.

YES!

10 lectures scheduled in **3** classrooms



Greedy algorithm

- **Idea.** Sort lectures in increasing order of **start time**: assign lecture to any **compatible** classroom.

IntPartition ($\{s_j, f_j\}$) // $r \leftarrow 0$ # of allocated rooms

1. Sort by **starting** time so that $s_1 \leq s_2 \leq \dots \leq s_n$

2. For $j = 1, \dots, n$

 If j compatible with some classroom k

 Schedule j in room k

 Else allocate new classroom $r + 1$

 Schedule j in room $r + 1$

$r \leftarrow r + 1$

} How to do it in $O(\log r)$

OBS. # rm needed \geq
depth of input intervals
(i.e., Max. number of
lectures that overlap)

- **Running time.** $O(n \log n)$
- **Optimality.** #Rm allocated = depth of input intervals

